GENERAL ASSEMBLY

DATA SCIENCE

SYD DAT 6

Week 4 – Clustering
Wednesday 2nd November

AGENDA

1. Motivation / Review
2. What is Clustering?
3. What is K-Means and how does it work?
4. Lab
5. Discussion


Atlassian

DATA SCIENCE PART TIME COURSE

WHAT IS CLUSTERING AND WHY DO IT?

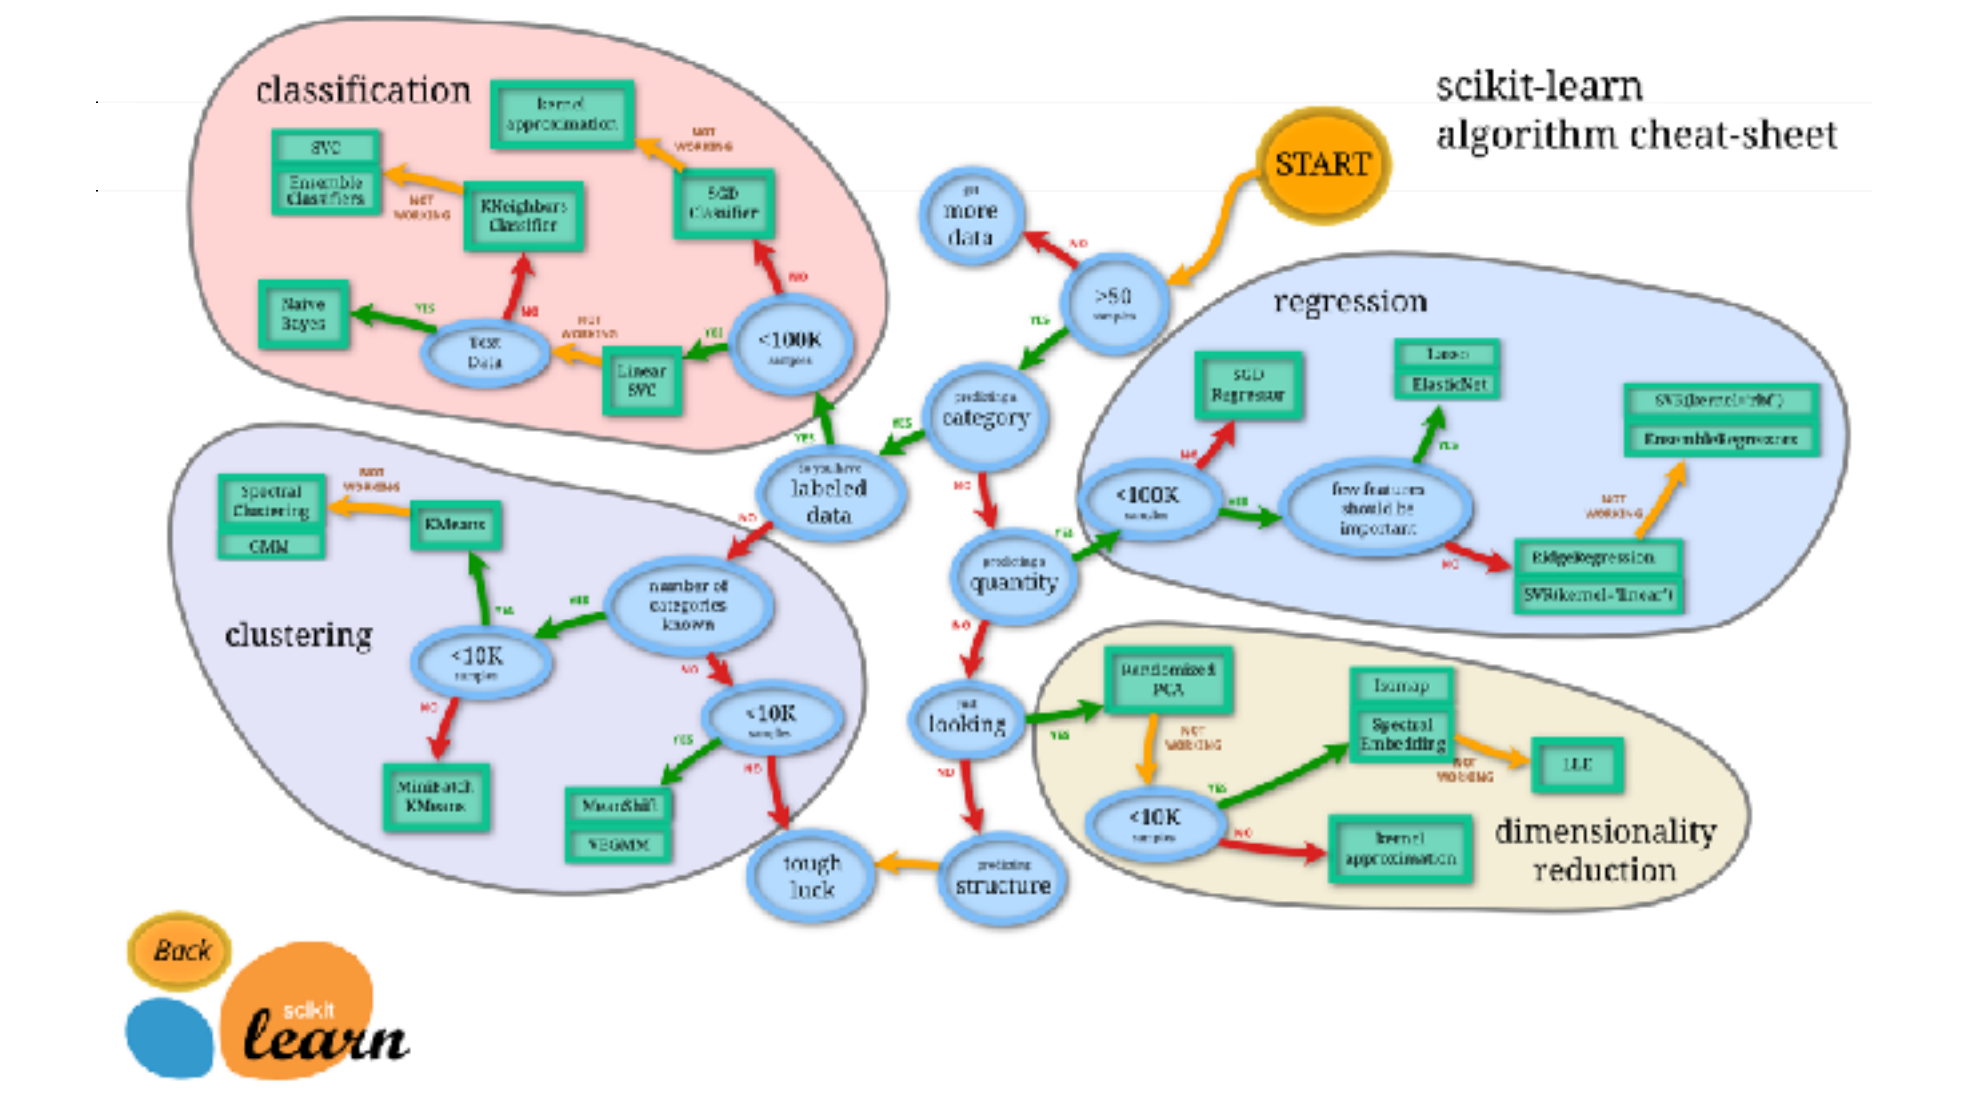
scikit-learn
algorithm cheat-sheet
START
classification
regression
clustering
dimensionality
reduction
more
data
>50
samples
<100K
samples
predicting
category
predicting
quantity
just
looking
predicting
structure
tough
luck
labeled
data
number of
categories
known
few features
should be
important
<10K
samples
<10K
samples
<10K
samples
<10K
samples
SVC
Ensemble
classifiers
KNeighbors
Classifier
SGD
Classifier
Kernel
approximation
Naive
Bayes
Text
Data
Linear
SVC
SGD
Regressor
Lasso
ElasticNet
SVR(kernel='rbf')
EnsembleRegressors
RidgeRegression
SVR(kernel='linear')
Spectral
Clustering
GMM
KMeans
MiniBatch
KMeans
MeanShift
VBGMM
Randomized
PCA
Isomap
Spectral
Embedding
LLE
Kernel
approximation
NOT
WORKING
YES
NO
Back
scikit
learn

MNIST

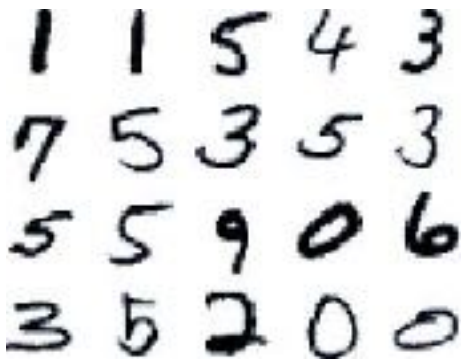

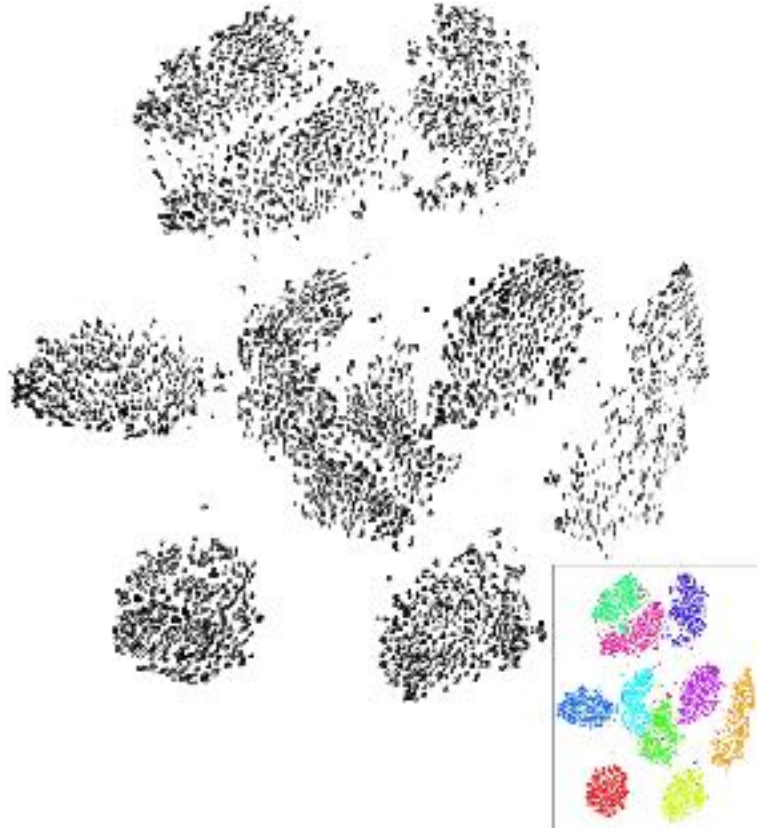

Olivetti Faces

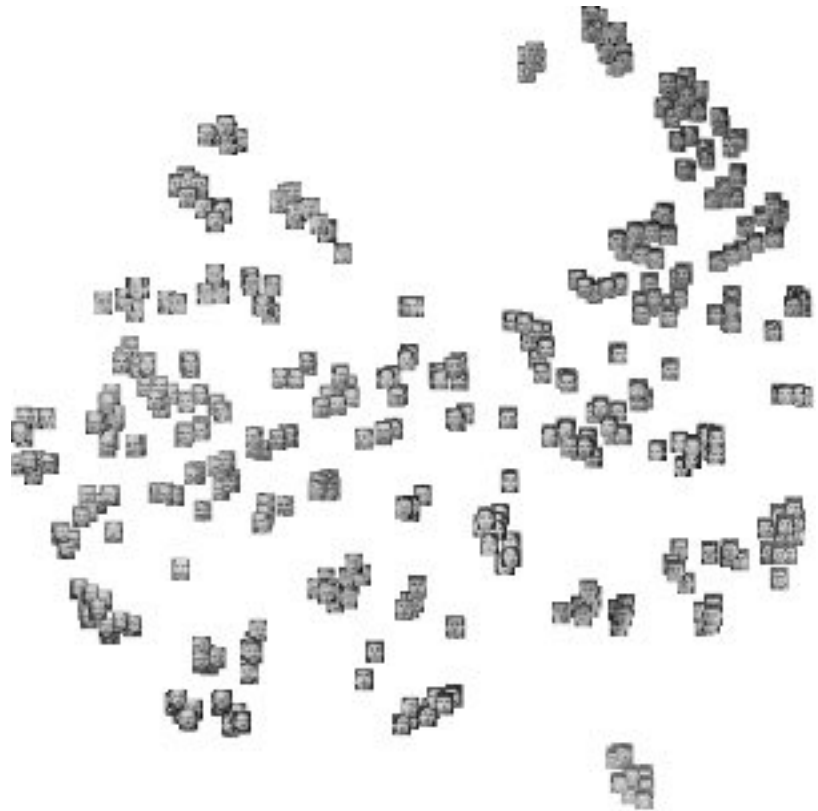

- What is a Cluster?
- Why would we do this?
- What is K-Means?

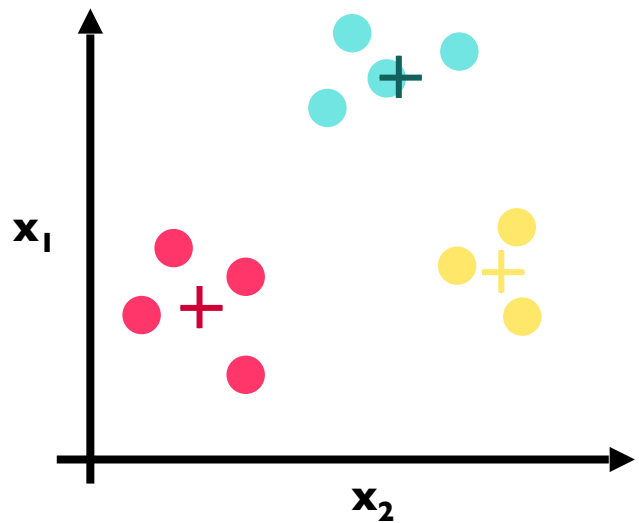

Recall unsupervised learning is when we are trying to find interesting patterns or groups in our data. We don't have a variable we are trying to predict (a Y value).

Clustering aims to discover subgroups in our data where the points are similar to each other. So we have a collection of groups and all points belonging to the same group are similar. Points in different groups are different to each other.

We have to decide what variables we will construct the groups on. What makes them different (or similar)?

To enhance our understanding of a dataset by dividing the data into groups.

Clustering provides a layer of abstraction from individual data points.

The goal is to extract and enhance the natural structure of the data

Marketing teams might want to group customers into like groups as a way of summarising the data

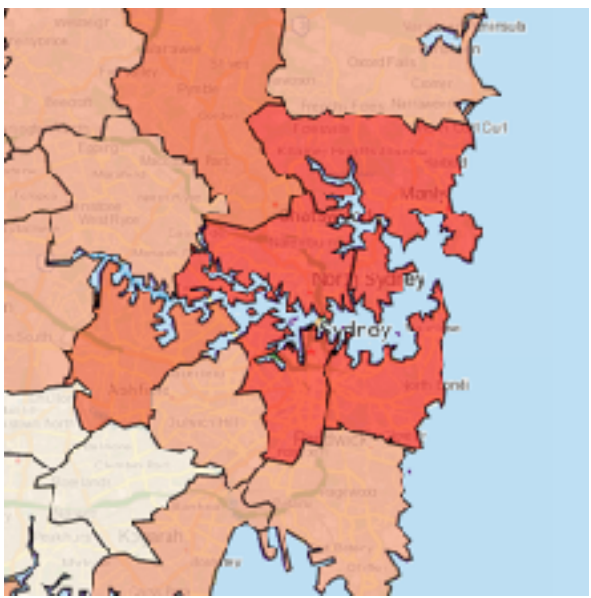

Financial groups may want to group transactions into like groups as a way to find unusual payments

Genetics data can be clustered to identify ancestry

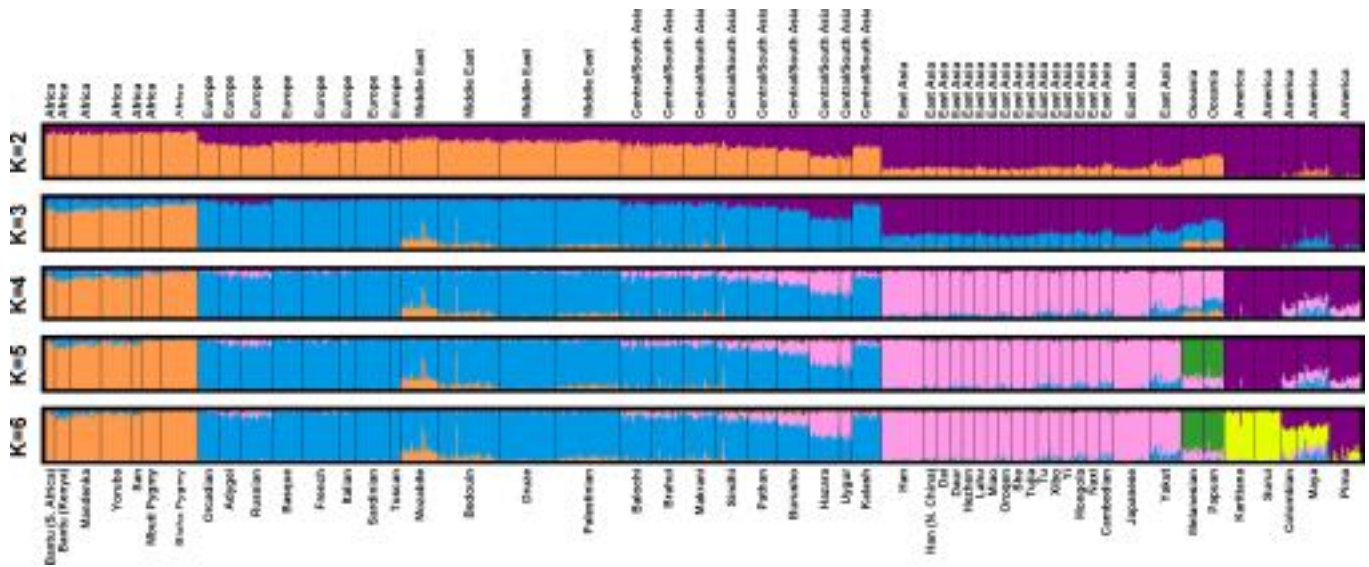

HOW DO WE CLUSTER DATA?

1) Choose k initial centroids (note that k is an input)

2) For each point:
   - find distance to each centroid
   - assign point to nearest centroid

3) Recalculate centroid positions

4) Repeat steps 2-3 until stopping criteria met

There are several options:

- randomly (but may yield divergent behavior)

- perform alternative clustering task, use resulting centroids as initial k-means centroids

- start with global centroid, choose point at max distance, repeat (but might select outlier)

The similarity criterion is determined by the measure we choose.

In the case of k-means clustering, the similarity metric is the **Euclidian distance:**

$$d(x_1, x_2) = \sqrt{\sum_{i=1}^{N} (x_{1i} - x_{2i})^2}$$

Q: How do we re-compute the positions of the centres at each iteration of the algorithm?

A: By calculating the centroid (i.e., the geometric centre)

We iterate until some stopping criteria are met; in general, suitable convergence is achieved in a small number of steps.

Stopping criteria can be based on the centroids (eg, if positions change by no more than ε) or on the points (eg, if no more than x% change clusters between iterations).

1) Choose k initial centroids

2) For each point:
   - find distance to each centroid
   - assign point to nearest centroid

3) Recalculate centroid positions

4) Repeat steps 2-3 until stopping criteria met

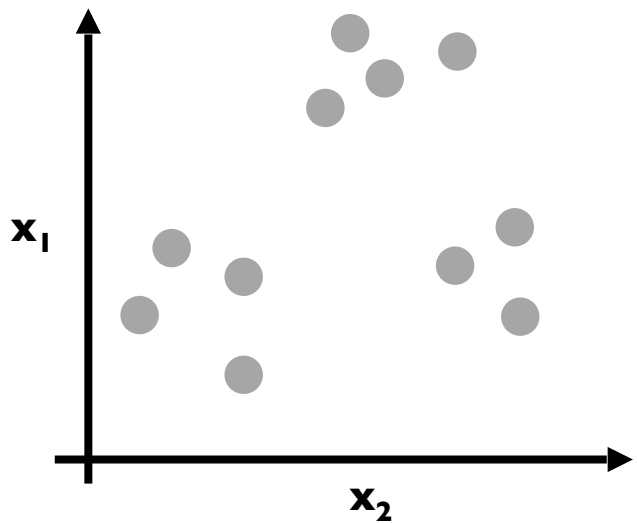

1) **Choose k initial centroids**

2) For each point:
   - find distance to each centroid
   - assign point to nearest centroid

3) Recalculate centroid positions

4) Repeat steps 2-3 until stopping criteria met

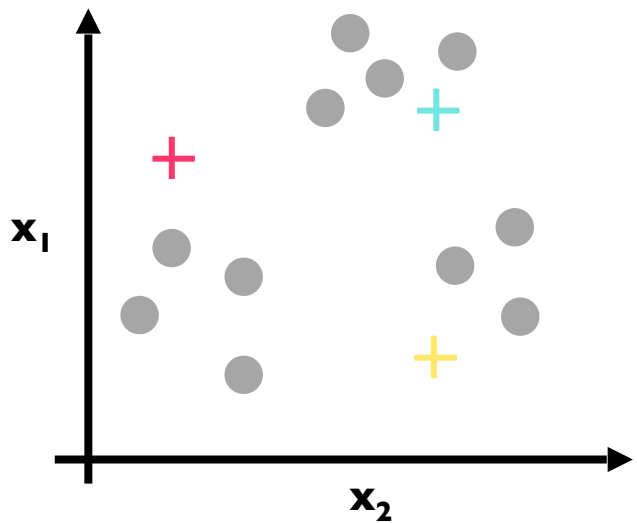

1) Choose k initial centroids

2) For each point:
   - **find distance to each centroid**
   - assign point to nearest centroid

3) Recalculate centroid positions

4) Repeat steps 2-3 until stopping criteria met

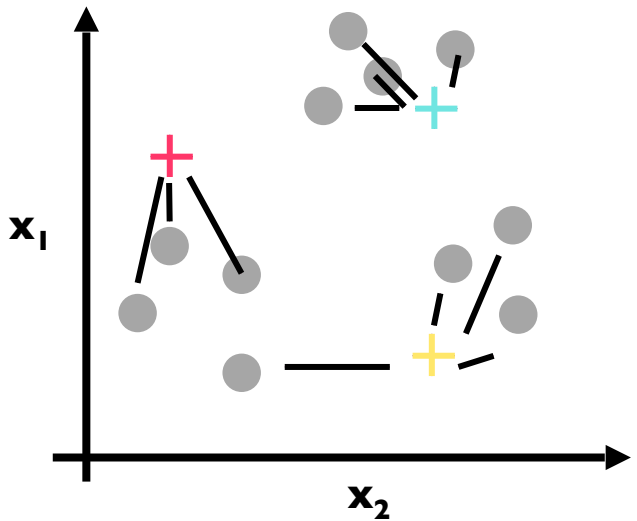

1) Choose k initial centroids

2) For each point:
   - find distance to each centroid
   - **assign point to nearest centroid**

3) Recalculate centroid positions

4) Repeat steps 2-3 until stopping criteria met

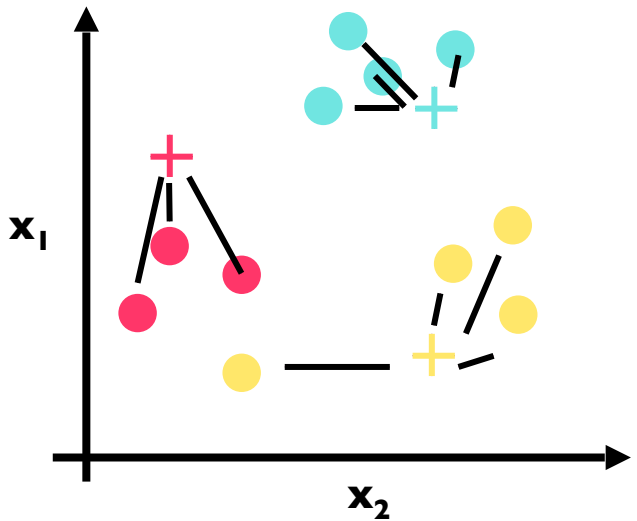

1) Choose k initial centroids

2) For each point:
   - find distance to each centroid
   - assign point to nearest centroid

3) Recalculate centroid positions

4) Repeat steps 2-3 until stopping criteria met

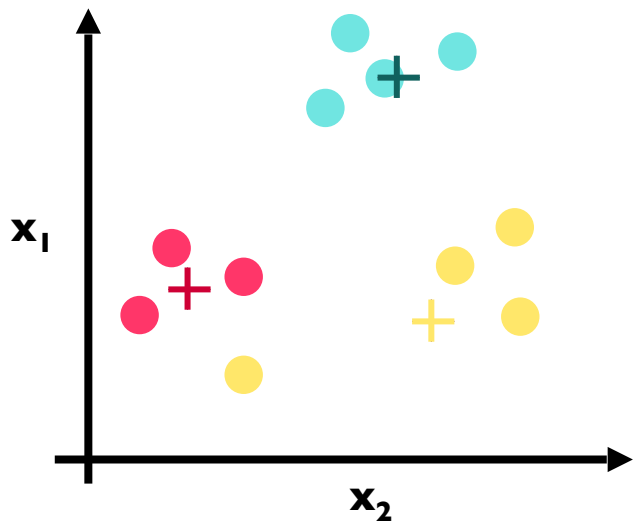

1) Choose k initial centroids

2) For each point:
   - find distance to each centroid
   - assign point to nearest centroid

3) Recalculate centroid positions

4) Repeat steps 2-3 until stopping criteria met

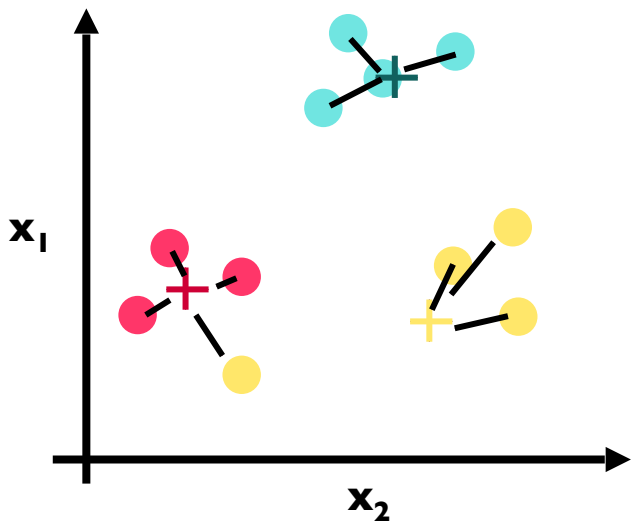

1) Choose k initial centroids

2) For each point:
   - find distance to each centroid
   - assign point to nearest centroid

3) Recalculate centroid positions

4) Repeat steps 2-3 until stopping criteria met

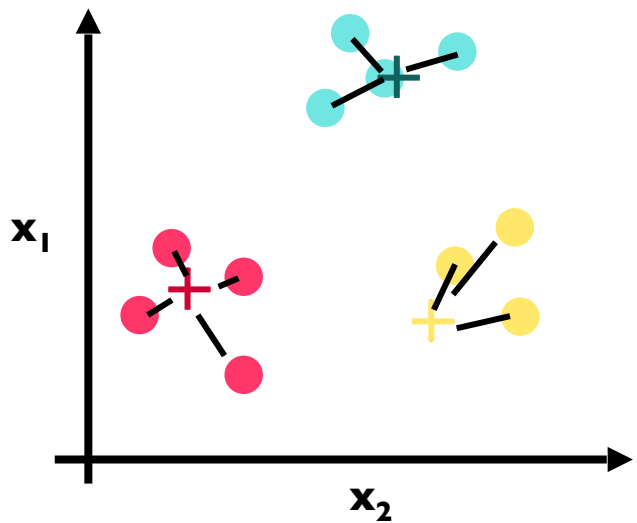

1) Choose k initial centroids

2) For each point:
   - find distance to each centroid
   - assign point to nearest centroid

3) Recalculate centroid positions

4) Repeat steps 2-3 until stopping criteria met

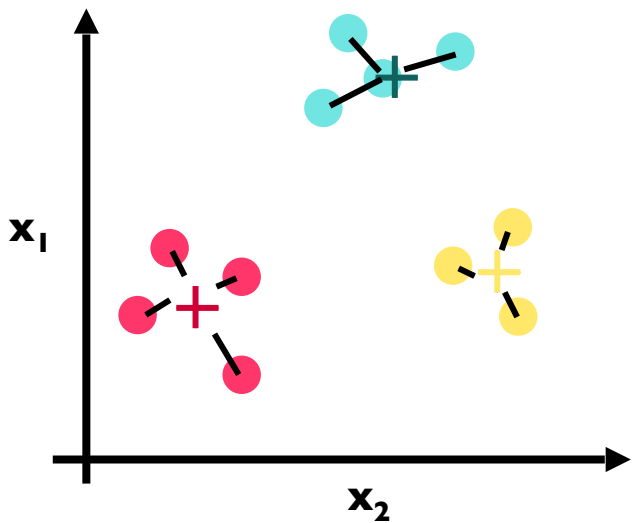

1) Choose k initial centroids

2) For each point:
   - find distance to each centroid
   - assign point to nearest centroid

3) Recalculate centroid positions

4) Repeat steps 2-3 until stopping criteria met

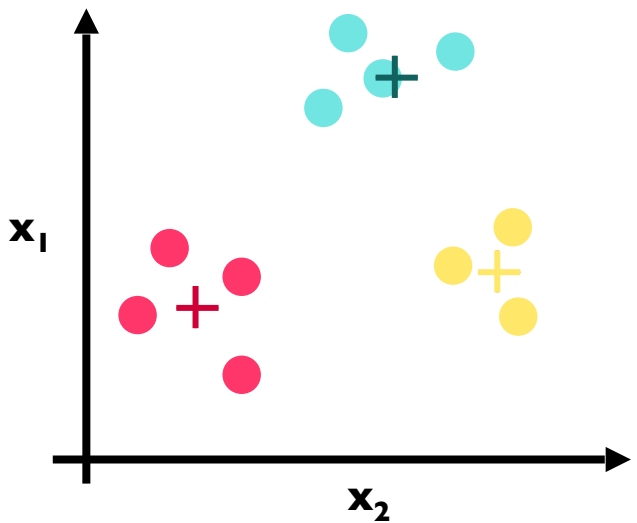

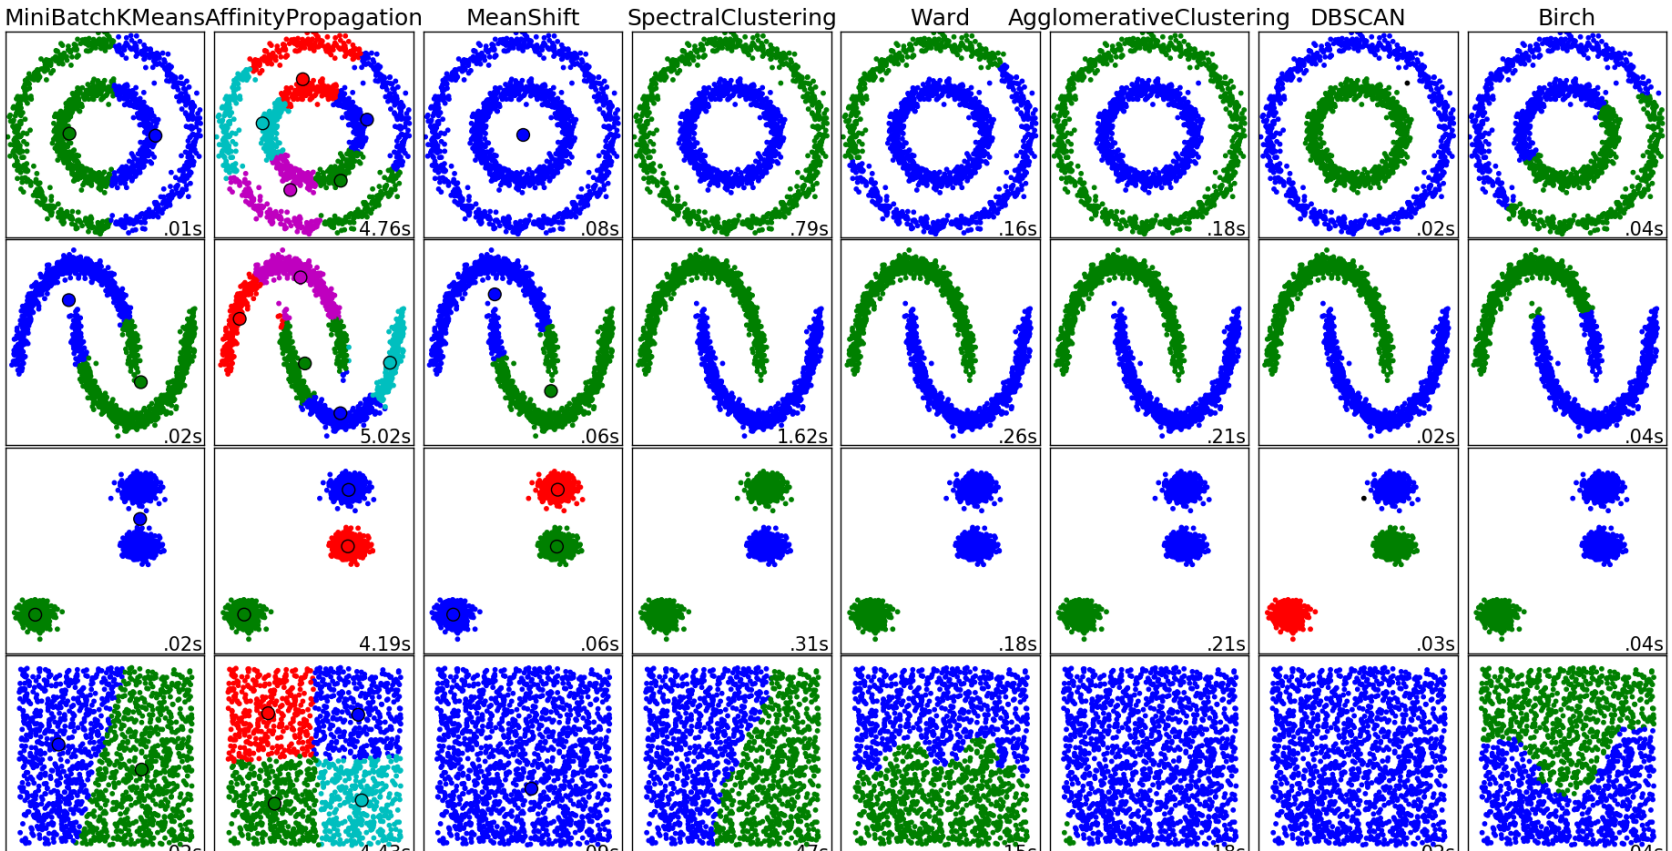
MiniBatchKMeans
AffinityPropagation
MeanShift
SpectralClustering
Ward
AgglomerativeClustering
DBSCAN
Birch
.01s
4.76s
.08s
.79s
.16s
.18s
.02s
.04s
.02s
5.02s
.06s
1.62s
.26s
.21s
.02s
.04s
.02s
4.19s
.06s
.31s
.18s
.21s
.03s
.04s
.02s
4.43s
.09s
.47s
.15s
.18s
.02s
.04s

DATA SCIENCE PART TIME COURSE

HOW DO WE KNOW OUR CLUSTERS ARE ANY GOOD?

In general, k-means will converge to a solution and return a partition of k clusters, even if no natural clusters exist in the data.

We will look at two validation metrics useful for partitional clustering, **cohesion** and **separation**.

Cohesion measures clustering effectiveness within a cluster.

$$\hat{C}(C_i) = \sum_{x \in C_i} d(x, c_i)$$

Separation measures clustering effectiveness between clusters.

$$\hat{S}(C_i, C_j) = d(c_i, c_j)$$

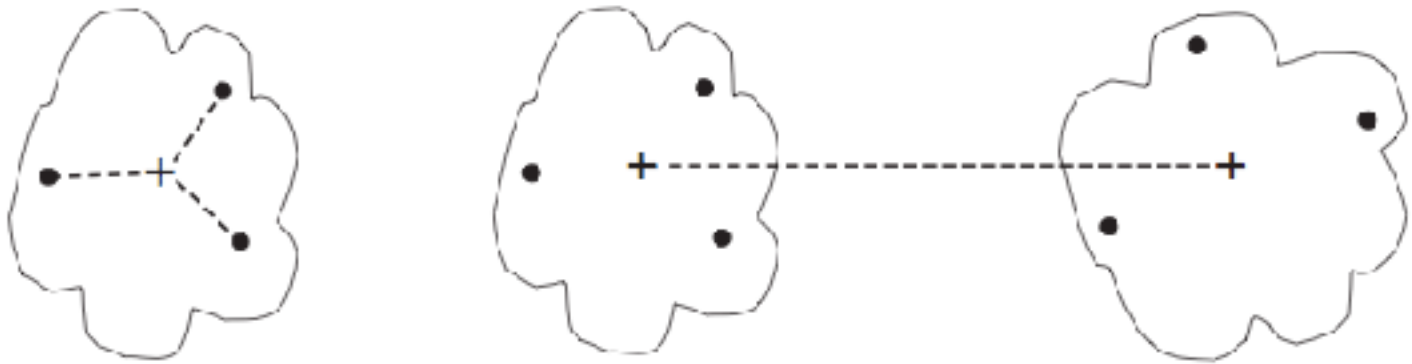

(a) Cohesion. (b) Separation.

Figure 8.28. Prototype-based view of cluster cohesion and separation.

source: http://www-users.cs.umn.edu/~kumar/dmbook/ch8.pdf

One useful measure than combines the ideas of cohesion and separation is the silhouette coefficient. For point x_i, this is given by:

$$SC_i = \frac{b_i - a_i}{max(a_i, b_i)}$$

such that:

a_i = average in-cluster distance to x_i

b_{ij} = average between-cluster distance to x_i

$b_i = \min_j(b_{ij})$

The silhouette coefficient can take values between -1 and 1.

In general, we want separation to be high and cohesion to be low. This corresponds to a value of SC close to +1.

A negative silhouette coefficient means the cluster radius is larger than the space between clusters, and thus clusters overlap

The silhouette coefficient for the cluster C_i is given by the average silhouette coefficient across all points in C_i:

$$SC(C_i) = \frac{1}{m_i} \sum_{x \in C_i} SC_i$$

The overall silhouette coefficient is given by the average silhouette coefficient across all clusters:

$$SC_{total} = \frac{1}{k} \sum_{1}^{k} SC(C_i)$$

One useful application of cluster validation is to determine the best number of clusters for your dataset.

Q: How would you do this?

A: By computing the SSE or SC for different values of k.

Ultimately, cluster validation and clustering in general are suggestive techniques that rely on human interpretation to be meaningful.

Art

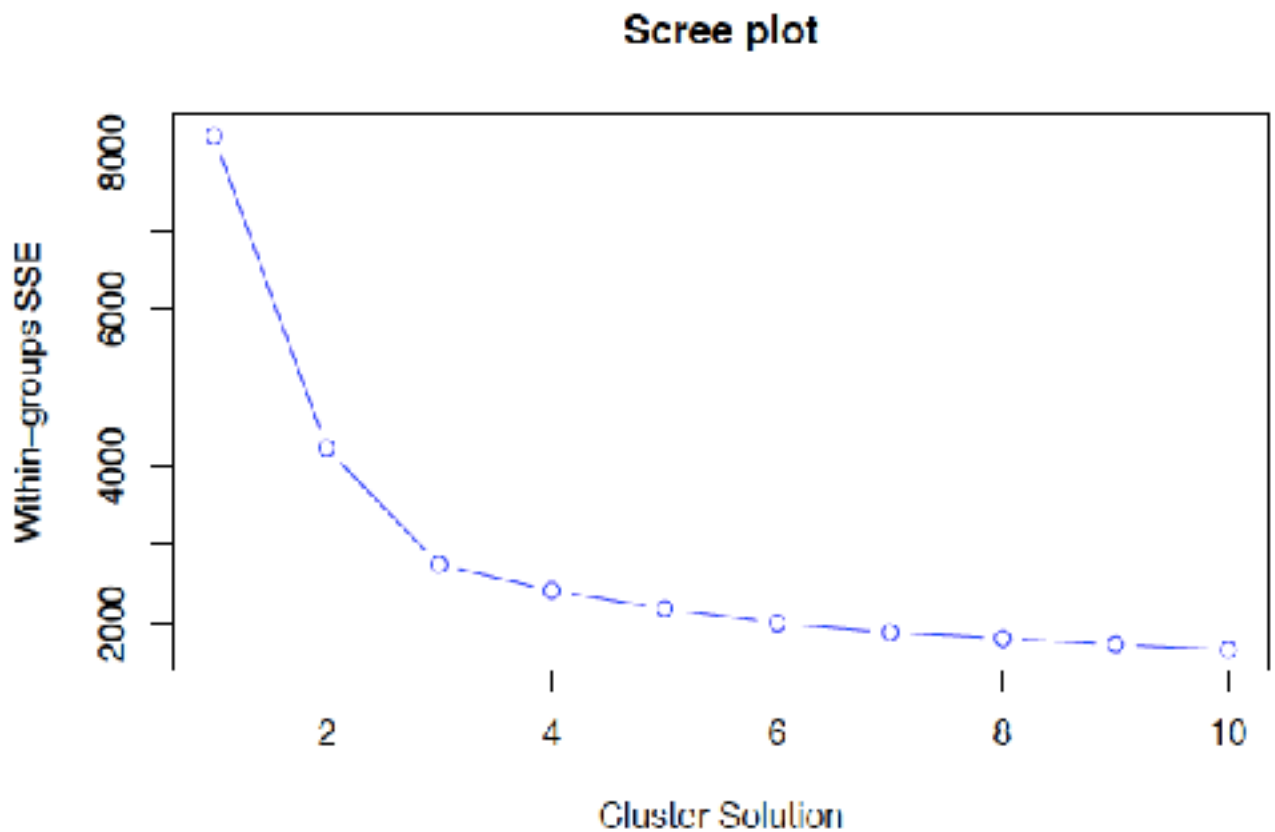
Scree plot
Within−groups SSE
2000
4000
6000
8000
2
4
6
8
10
Cluster Solution

Strengths:

K-means is a popular algorithm because of its computational efficiency and simple and intuitive nature.

Weaknesses:

However, K-means is highly scale dependent, and is not suitable for data with widely varying shapes and densities.

LAB

SYNCHING YOUR FORK WITH THE COURSE REPO

1. re-name your labs with lab_name.<yourname>.ipynb (to prevent a conflict)
2. cd <path to the root of your SYD_DAT_6 local repo>
3. commit your changes ahead of sync
   - git status
   - git add .
   - git commit -m "descriptive label for the commit"
   - git status
4. download new material from official course repo (upstream) and merge it
   - git checkout master (ensures you are in the master branch)
   - git fetch upstream
   - git merge upstream/master

HOMEWORK

Homework

- Homework 2 – Due Monday 14th of November

Read the following

- Chapter 10.3 of Introduction to Statistical Learning – Clustering Methods in Introduction to Statistical Learning (15 pages)
- Chapter 8.1 of Introduction to Statistical Learning – The Basics of Decision Trees